

# Machine Learning 520

## Advanced Machine Learning

---

### Lesson 1: Introduction to Advanced Machine Learning



## Today's Agenda

---

- Class Introduction
- Class policy and expectations
- How to succeed in this course
- Assignments, quizzes and milestones
- Participation and grading
- Grading expectations and supplementary material
- Introduction to Advanced Machine Learning
- Challenges in Machine Learning and how to overcome them



## Learning Objectives

---

By the end of this lesson, you should be able to

- Implement two strategies for multiclass classification (lab).
- Apply implementation to data set .
- Apply insights gained from exploring a real-world case to thinking about how to address real-world problems, starting with a machine learning problem, then applying techniques.



## Course Overview

---

- Lesson 01: Introduction to Advanced Machine Learning
- Lesson 02: Decision Tree
- Lesson 03: Random Forest & Gradient Boosted Trees
- Lesson 04: Support Vector Machine
- Lesson 05: Stacking and Blending
- Lesson 06: Unsupervised Learning
- Lesson 07: Natural Language Processing
- Lesson 08: Recommendation Systems
- Lesson 09: Forecasting
- Lesson 10: Building Machine Learning Applications



## Student Introduction

---

- Name and current role
- Why enrolled in data science certificate
- Course and program expectations



## Course Format

---

- On average, we spend about 40% on lecture and 60% on hands-on
- Let's have frequent discussions during the session and share diverse perspectives
- Let's take frequent short breaks to fit online format



## How to succeed in this course

---

- Pay attention during class and ask questions
- Run notebooks alongside but don't let it distract you too much
- Learn from peers by participating in discussion boards and by sharing your code after assignment due dates
- Do your best not to fall behind, because lessons build on each other
- Basic to intermediate Python programming concepts should be second nature by now



## **Assignments, quizzes and milestones**

---

- Assignments are due by 11:59 PM one week after lecture
- I will make no exceptions about assignment due dates
- Questions about the assignments should be posted on the discussion board on Canvas, where instructional team will answer them (or students if they wish)





## How to submit a great assignment

---

- Your code should run from start to finish, so once you finish the assignment, restart the notebook and run again to make sure nothing breaks
- Each step in the assignment should include (1) one code cell with a few comments in the code to point out key parts, and (2) a Markdown cell explaining your reasoning and (3) a Markdown cell explaining your conclusions
- Someone who doesn't know Python could still be able to read through and understand the analysis
- When faced with ambiguity, you are free to make assumptions as long as you (1) clearly state your assumptions and (2) it is a reasonable assumption





**let's take a tour through Canvas**



## Grading expectations

---

- Please use the discussion boards for questions on assignment
- DO NOT post your assignment code in discussion boards
- When necessary, it's generally OK to make an assumption about the assignment reqs - as long as you state your assumption
- More important to be on time than perfect:
- Meet the requirements with working code
- Write good comments for your code
- Write good explanations showing your line of thinking



## Participation and grading

---

activity	what you need to do	grade
Class and discussion board participation	be active in <b>discussion boards</b>	15%
Quizzes		10%
Weekly Lab Assignments	submit by <b>11:59 PM the week after</b>	75%



## Coding Environment

---

- We will be using browser-based Jupyter notebooks as our python environment
- Basics of Jupyter notebooks:
  - Code cells and Markdown cells
  - Running and re-running code cells - order matters!





**Break time**



# What Is Machine Learning?

---

- An algorithm is a self-contained set of **rules or instructions** used to *solve problems*
- **Machine learning** is the field of study that gives computers the ability to learn *without being explicitly programmed*, by using data to learn
- The *problems* ML algorithms try to solve are usually
  - prediction: **supervised learning**
  - finding structure in data: **unsupervised learning**
  - Mimicking human behavior: **reinforcement learning** (not covered)



## **Machine learning algorithm**

---

A machine learning algorithm walks into a bar.

The bartender asks, "What'll you have?"

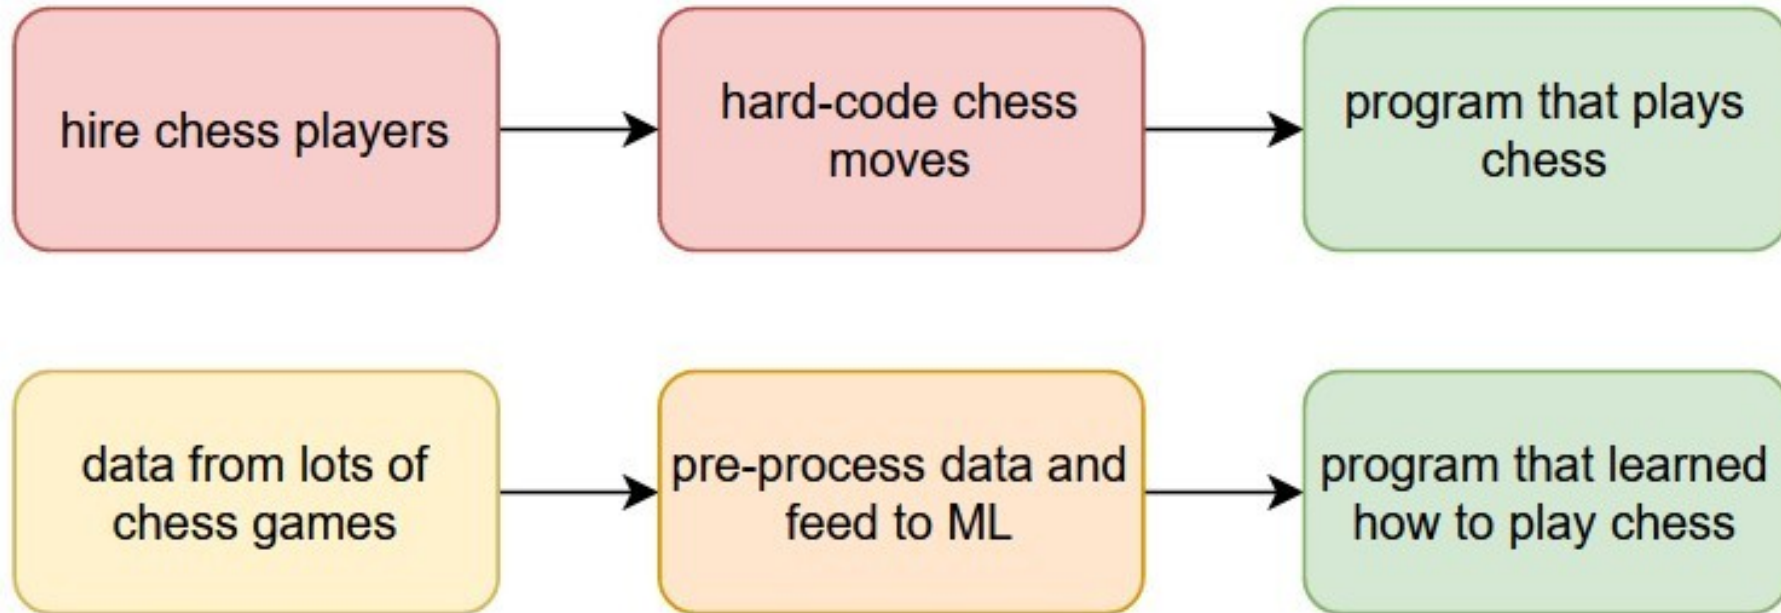
**The algorithm says, "What's everyone else having?"**






## Rule-based Vs Machine Learning?

---



# Types of Machine Learning Algorithms



## Supervised Learning

---

- Also called **predictive modeling** or **inference**
- Look at some current examples (**labeled data**) and find a model that can predict future examples (**unlabeled data**)
- The **target variable** or **label** is what we want to predict
  - **Regression** algorithms are used with a **numeric target**
  - **Classification** algorithms are used with a **categorical target**
- By comparing predictions with the actual labels, we can evaluate our model's accuracy (hence the term *supervised*)



## Unsupervised Learning

---

- Also called data-mining / pattern recognition / structure discovery
- Look at **unlabeled data** and find general patterns
- More subjective and difficult to evaluate and interpret, and hence it is far **less common** than supervised learning
- **Clustering** is the most common example
  - K-means clustering
  - Variable clustering / dimensionality reduction
  - Word clouds (kind of)



# Reinforcement Learning

---

- Learn to take actions in an environment to maximize expected future



## Discussion


Task	Learning Type
Housing price estimates on Redfin	?
Product recommendations on Amazon	?
Voice recognition on Alexa	?
Play the game of Go	?
Grouping Air Bnb listings into neighborhoods	?
Predict delivery times for Uber Eats	?
Segmenting US population for targeted advertising	?
Learn to play Starcraft	?
File a tax return	?





## Multi-Class Classification


- Multi-class classification generalizes the problem where classifier
  - is expected to distinguish between more than two classes
  - Hand-written digits (0, 1, ..., 9)
  - Image classification (cat, dog, human etc.)
  - Classifying different species of flowers



W

# Multi-Class Classification – One vs. All

- One vs. All is a way to handle multi-class classification
- Training
  - Fit a binary classifier per class, with the samples of that class as positive samples and all other samples as negative
  - Base classifiers produce a real-valued confidence score for its decision, instead of just a class label
- Prediction
  - Apply all trained classifiers to the new unseen data point and select the class with highest score





## Multi-Class Classification – One vs. All


- A second strategy is One vs. One

- **Training**

- Fit a binary classifier for every pair of classes. Therefore, if there are  $N$  classes, we need  $N*(N-1)/2$  classifiers

- **Prediction**

- A voting scheme is applied
  - All  $N*(N-1)/2$  classifiers are applied to an unseen sample and the class that gets the highest number of votes is selected



## Underfitting vs. Overfitting

---

- Fitting means **learning** when we call the `.fit()` method
- If a model performs poorly on the training data, then it almost certainly will perform poorly on the test data as well: we say the model is **underfitting** (not learning enough)
- If a model performs well on the training data, but poorly on the test data: we say the model is **overfitting** (it's learning the signal but also "learning" the noise in the training data, and hence fails to generalize)
- A good model is one that neither underfits nor overfits



## Overfitting and Complexity


---

- More simple models tend to underfit, because they are more likely to oversimplify (not pick up enough signal)
  - Analogy: models can be prejudiced too, we call it **bias**
- Complex models tend to overfit, because they are so eager to pick up any signal that they also grab noise disguised as signal
  - **Analogy:** people who **read too much into** something
- The trick is to find the happy medium

*Everything should be made as simple as possible, but not simpler. -Albert Einstein*

(also look up **Occam's Razor**)





## What Are Hyper-Parameters?

---

- Almost all algorithm have ways they can be "tuned" through different **hyper-parameter** choices
- Some hyper-parameters are very generic, e.g. **learning rate**
- Most hyper-parameters are algorithm-specific, such as
  - For tree-based algorithms: tree depth, min leaf size
- ML algorithms **cannot** directly learn optimal hyper-parameter values during training
- If we don't specify them, they usually default to "reasonable" values



## Model Selection and Validation

---

- So how do we know what hyper-parameter values to pick?
    - Mostly through trial and error, also called **model selection**
  - If we want to tune our hyper-parameters, we also need **validation data** in addition to training and testing
1. Train many models, each with different hyper-parameters, and evaluate their performance on the validation data
  2. The model that performs best on the validation data wins
  3. Check how the winner model performs on the test data



## Regularization

---

- **Reduces overfitting** and the extent of it can be adjusted or **tuned**
- Instead of minimizing prediction error only, minimize **prediction error + a penalty term** where the penalty term is higher when model coefficients are higher
- Prefers models with **smaller coefficients** (features must be normalized), **unless** they significantly improve prediction



# L1 vs L2 Norm Regularization

- has sharp edges
- has round edges
- can drop features




Image source: Elements of Statistical Learning



## Types of Regularization

---

- **L2 regularization** (such as in **ridge regression**) is good when we have fewer and more relevant features and we want to keep all
- **L1 regularization** (such as in **lasso regression**) is good when we have lot's of features and many are probably not relevant
- We can also combine **L1 regularization** and **L2 regularization**, such as in **elastic-net regression**, where we tune  $\lambda_1$  in addition to  $\lambda_2$



## Test Data vs. Validation Data

---

So why not combine test data and validation data?

- Because the test data is used **once** with the **final model** to get an **unbiased estimate** of performance (prediction error)
- The validation data is used **many times** so we can compare the performance of models trained with different sets of hyper-parameters, a.k.a. **hyper-parameter tuning**
- If we also use the test data to tune hyper-parameters, we are over-using it and its estimate of performance will not be so unbiased anymore



# Cross Validation

- Choose number of folds
- Train k times, each using folds to train and the th to evaluate
- Aggregate the evaluation results from all models into one measure

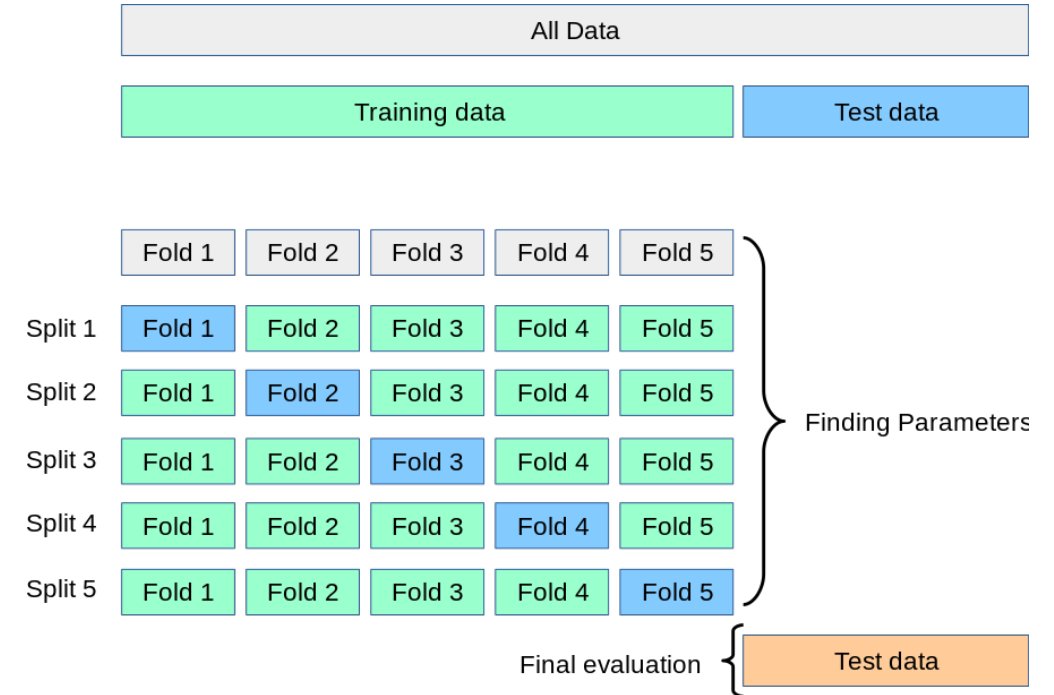
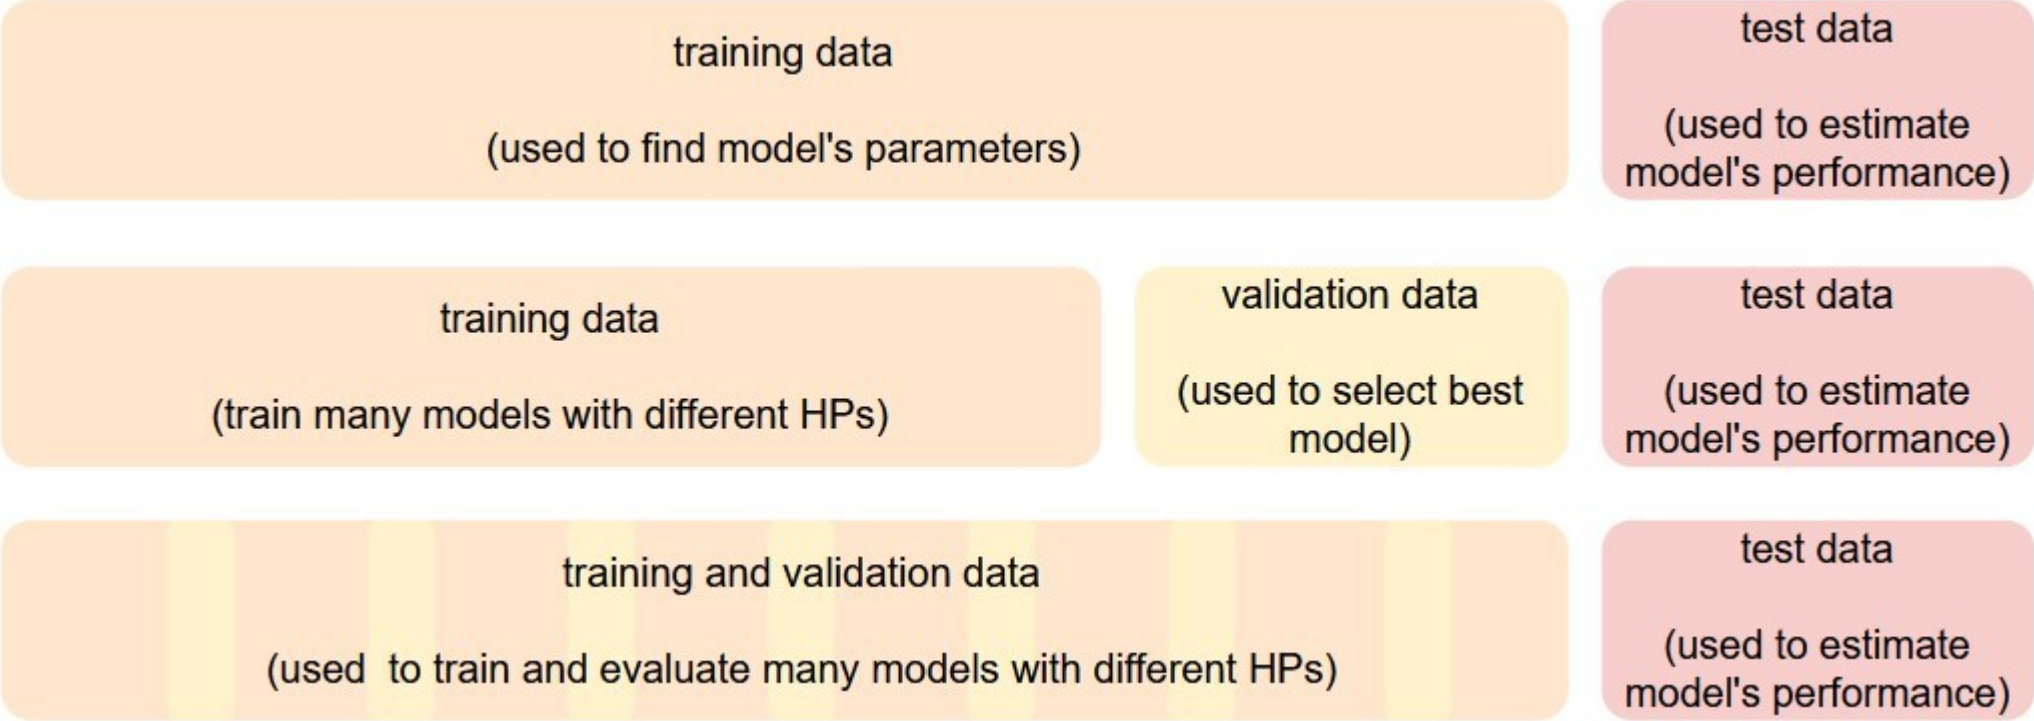



Image source: sciit-learn.org





## The Confusion Matrix

It's really not that confusing!

	predicted positive	predicted negative
<b>actually positive</b>	true positive TP	false negative FN
<b>actually negative</b>	false positive FP	true negative TN

For TP/FP/TN/FN

- The second letter indicates what the prediction was
- The first letter indicates if the prediction was right or not



## Discussion


---

- We saw there a binary classification model can make two kinds of errors: FP and FN
- For the following scenarios, say what kind of error is more costly (use common sense)
  - **Credit card fraud detection**: someone impersonates you to use your credit card
  - **Medical diagnosis**: finding out who has a disease
  - **Information retrieval**: finding relevant web pages based on a search query



# Precision and Recall

- **Precision** is the percentage of positive predictions that were actually positive
- **Recall** is the percentage of positive cases that we correctly predicted as such



How many selected items are relevant?

Precision =



How many relevant items are selected?

Recall =





## Accuracy, Precision and Recall

---

- For rare events usually TN far exceeds TP, FN, or FP, inflating accuracy, but precision and recall don't have TN in it



## Discussion

---

Here's an analogy that shows why we should evaluate a classification model's accuracy using **both** precision and recall:

- When you stand witness in a court of law, you are asked to tell the truth:
  - **The whole truth**: no lie of omission
  - **Nothing but the truth**: no lies
- Relate the above two statements to precision and recall





**Notebook time**





**The end**

